



# Chapter 1

## Introduction

**Da-Wei Chang**

**CSIE.NCKU**

Source: Abraham Silberschatz, Peter B. Galvin, and Greg Gagne, "Operating System Concepts", 10th Edition, Wiley.

# Objectives



- Introduce the major operating system components and functions
- To provide coverage of basic computer system organization
- Many of the topics will be discussed in detail in the following chapters

# Outline

- What Is an Operating System
- Computer-System Organization/Architecture
- Operating-System Structure/Operations
- Process Management
- Memory Management
- Storage Management
- Protection and Security
- Virtualization
- Basic Kernel Data Structures
- Computing Environments
- Open-Source Operating Systems

# What is an Operating System?

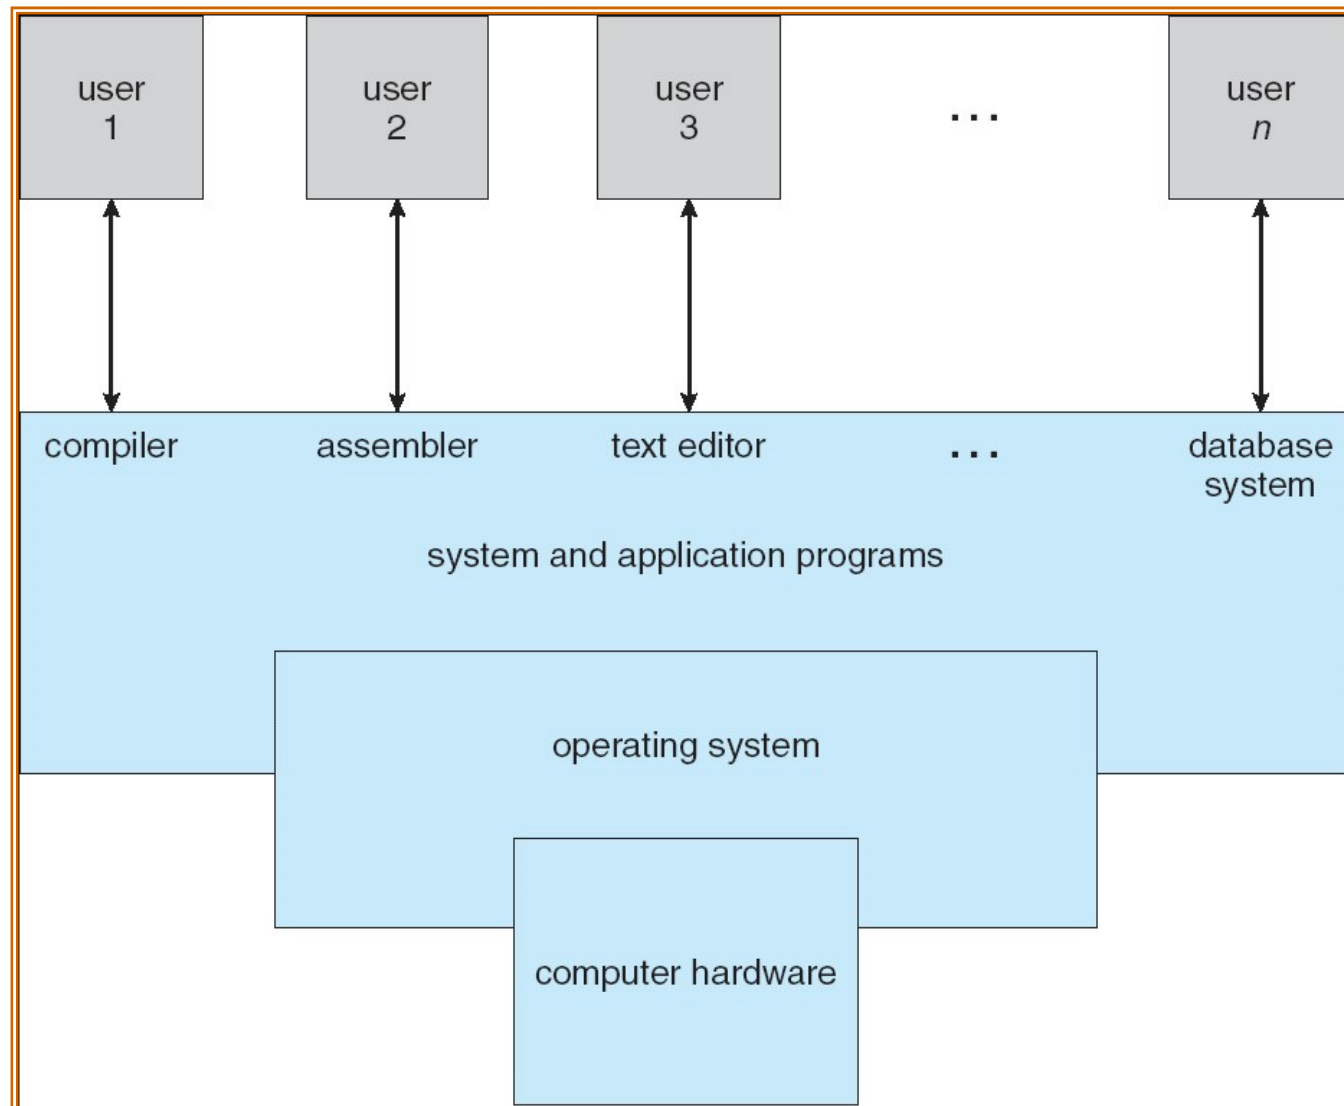
- A **program** that acts as an intermediary between **user programs/applications** and the **computer hardware**
- Operating system goals
  - Execute user programs
  - Make the computer system convenient to use
    - Applications do not have to deal with HW directly
    - OS ensures **efficient** resource **sharing**

# Computer System Structure



- A computer system can be divided into 4 components (from bottom to top)
  - **Hardware** – provides basic computing resources
    - CPU, memory, I/O devices
  - **Operating system**
    - Controls and coordinates use of hardware among various applications and users
  - **System and application programs** – solve the computing problems of the users
    - Word processors, compilers, web browsers, database systems, video games
  - **Users**
    - People, machines, other computers

# Four Components of a Computer System



# Operating System Definition

- **What is an operating system?**
- No universally accepted definitions
- “Everything a vendor ships when you order an operating system” is good approximation
  - But varies widely
- “The one program running at all times on the computer” (i.e., the **kernel**) Everything else is either a **system program** (ships with the operating system) or an **application program**

# Operating System Definition

- Two points of view (in the text book)
  - OS is a **resource allocator**
    - Manages all resources
    - Applications do not have to manage resources directly
      - OS hides HW interface
    - Decides between conflicting requests for **efficient** and **fair** resource use
  - OS is a **control program**
    - Controls execution of programs to prevent **errors** and **improper use** of the computer
- Other definitions
  - A program that can monitor the other programs
  - A program that that allows the other programs to **share** resources in a **controlled** manner
    - CPU, memory, IO devices....

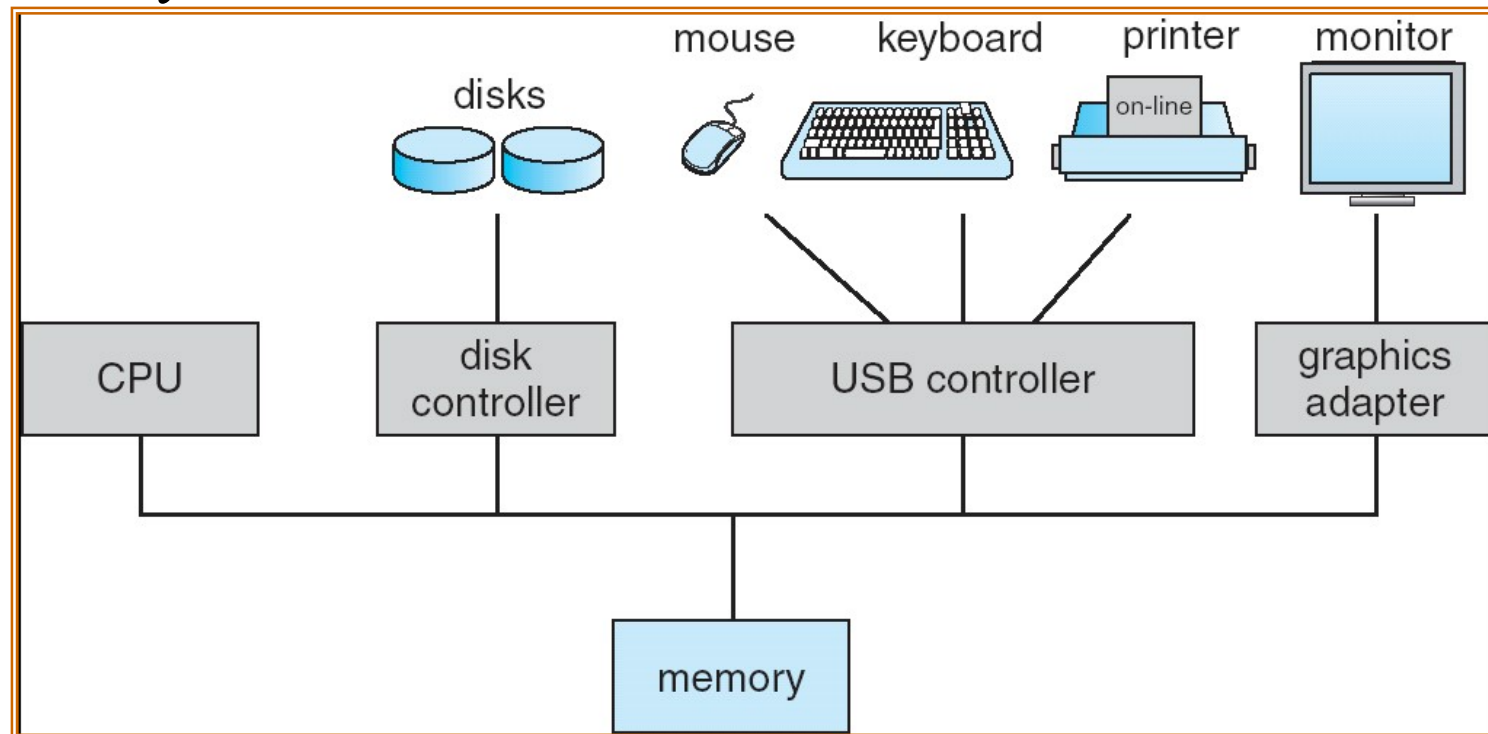


# Computer Startup

- **Who loads the OS ?**
- **Bootstrap program** is loaded at power-up or reboot
  - Typically stored in ROM or EEPROM, generally known as **firmware**
  - Jobs
    - Initializes all aspects of system
    - Loads operating system kernel and starts OS execution
- A bootstrap program is also called a **boot loader**
- Who loads the boot loader?
  - In PC, BIOS loads the (in-disk) boot loader
  - In many other platforms, boot loader is placed at a predefined memory address (in ROM or EEPROM)
    - Hardware jumps to the boot loader directly

# Computer System Organization

- Computer-system operation
  - One or more CPUs, device controllers connect through common bus providing access to shared memory
  - Concurrent execution of CPUs and devices competing for bus cycles



# Computer-System Operation


- I/O devices and the CPU can execute concurrently.
- Each device controller is in charge of a particular device type.
- Each device controller has a local buffer.
- Writing data to device a part of OS
  - (Device driver running on the) CPU moves data from main memory to local buffers
  - Device controller writes the data from the local buffer to the device
  - Device controller informs CPU that it has finished its operation by generating an interrupt
- One of the jobs of an operating system is to manage interrupts.

# Common Functions of Interrupts

- Interrupt transfers control to the **interrupt service routine (ISR)**, generally through the *interrupt vector (table)*
  - *interrupt vector* contains the **addresses or instructions** of all the service routines
    - X86: addresses
    - ARM: instructions
- The processor must save the address of the **interrupted** instruction before jumping to the ISR
- Generally, interrupts are raised by IO device controllers
  - **Hardware** interrupts

# Common Functions of Interrupts

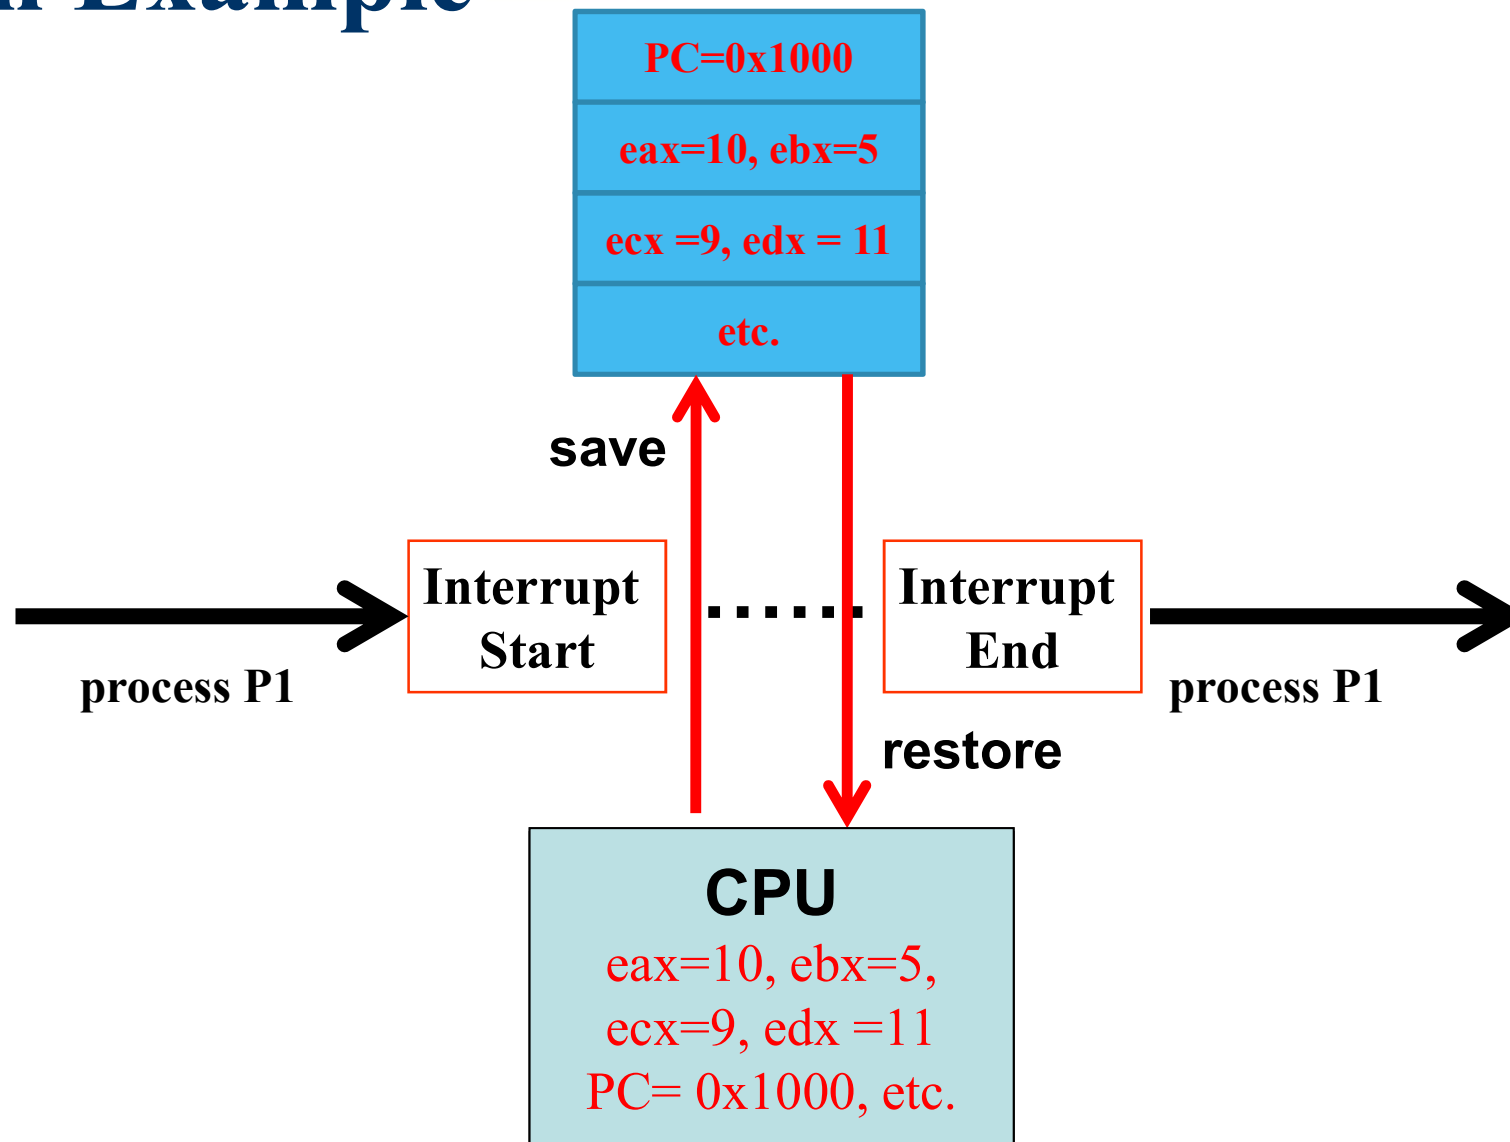
- However, *software* can also generate/raise interrupts
  - Software interrupts (or *traps*)
- A *trap* is a software-generated interrupt caused either by an *error* or a *user request*
  - Error
    - Division by zero, invalid memory access...
  - User request
    - Requests for OS services
      - *syscall/int* instruction in x86
      - *svc/swi* instruction in ARM
- An operating system is *interrupt driven*!



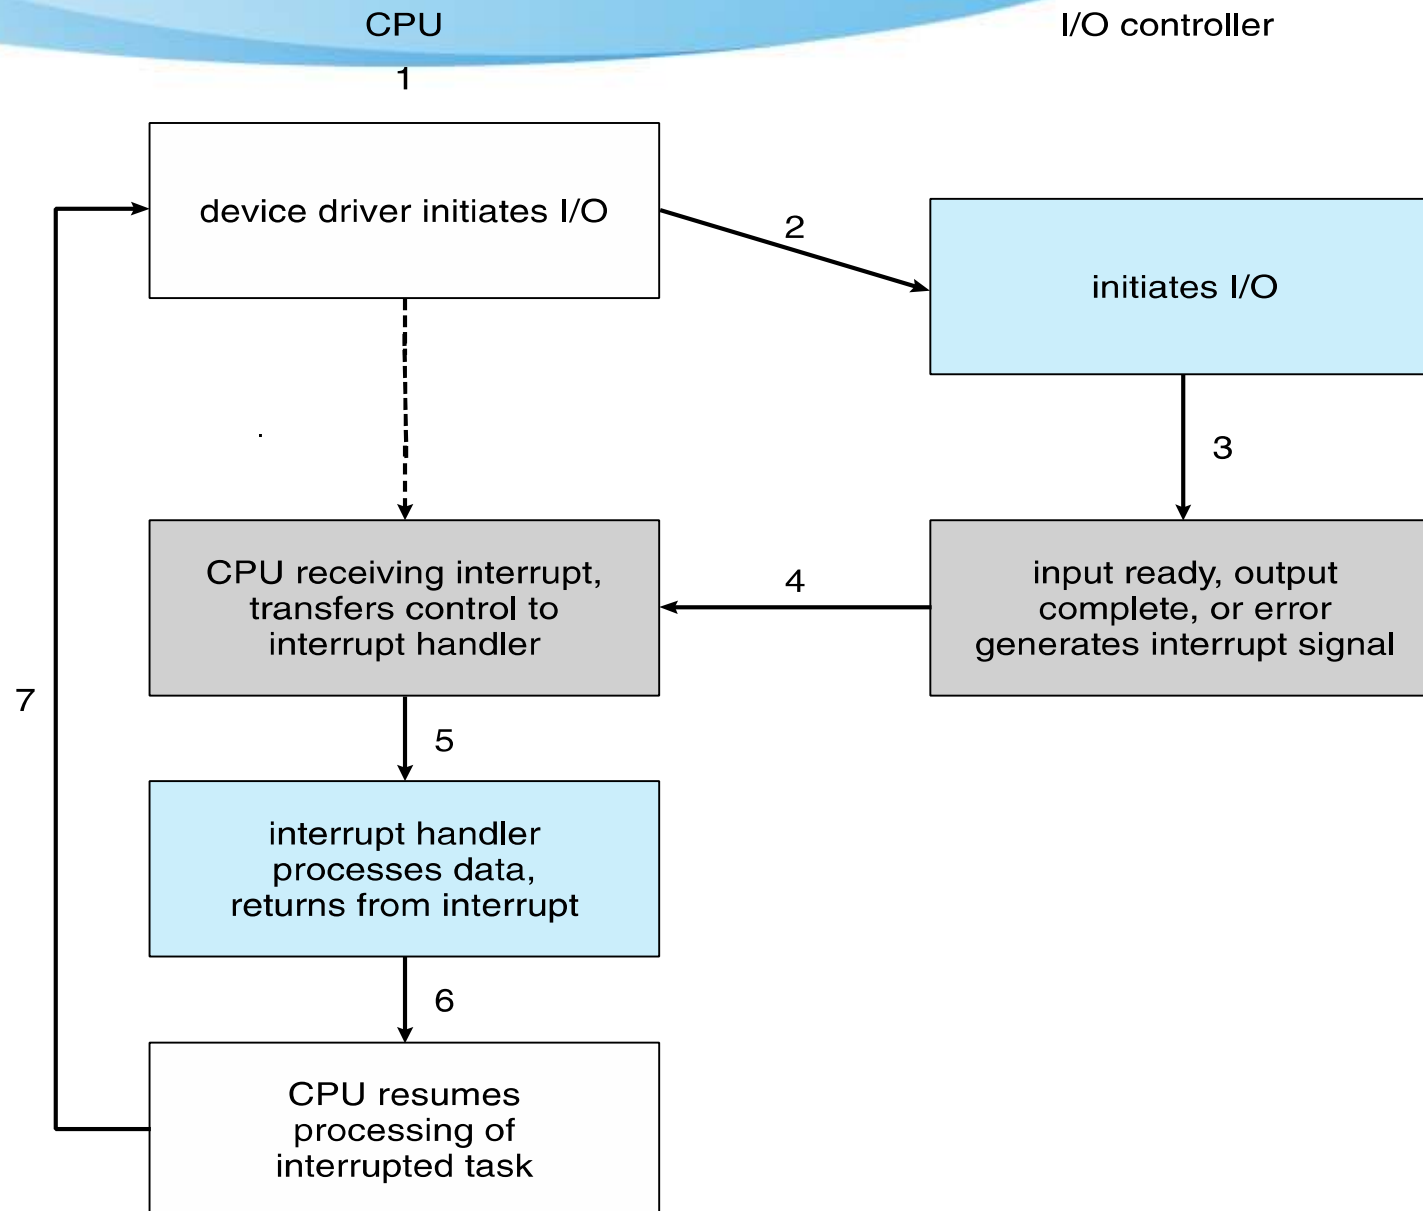
# Interrupt Handling

- Save the CPU state (*registers*, including program counter)
  - Done by HW or OS
- Locate the corresponding Interrupt Service Routine (ISR) or interrupt handler
  - How do you know the address of the ISR?
    - Interrupt vector table (IVT)
- Execute the ISR
- Continue the interrupted work by restoring the CPU state

# Saving/Restoring CPU State – An Example




# Interrupt-drive I/O Cycle





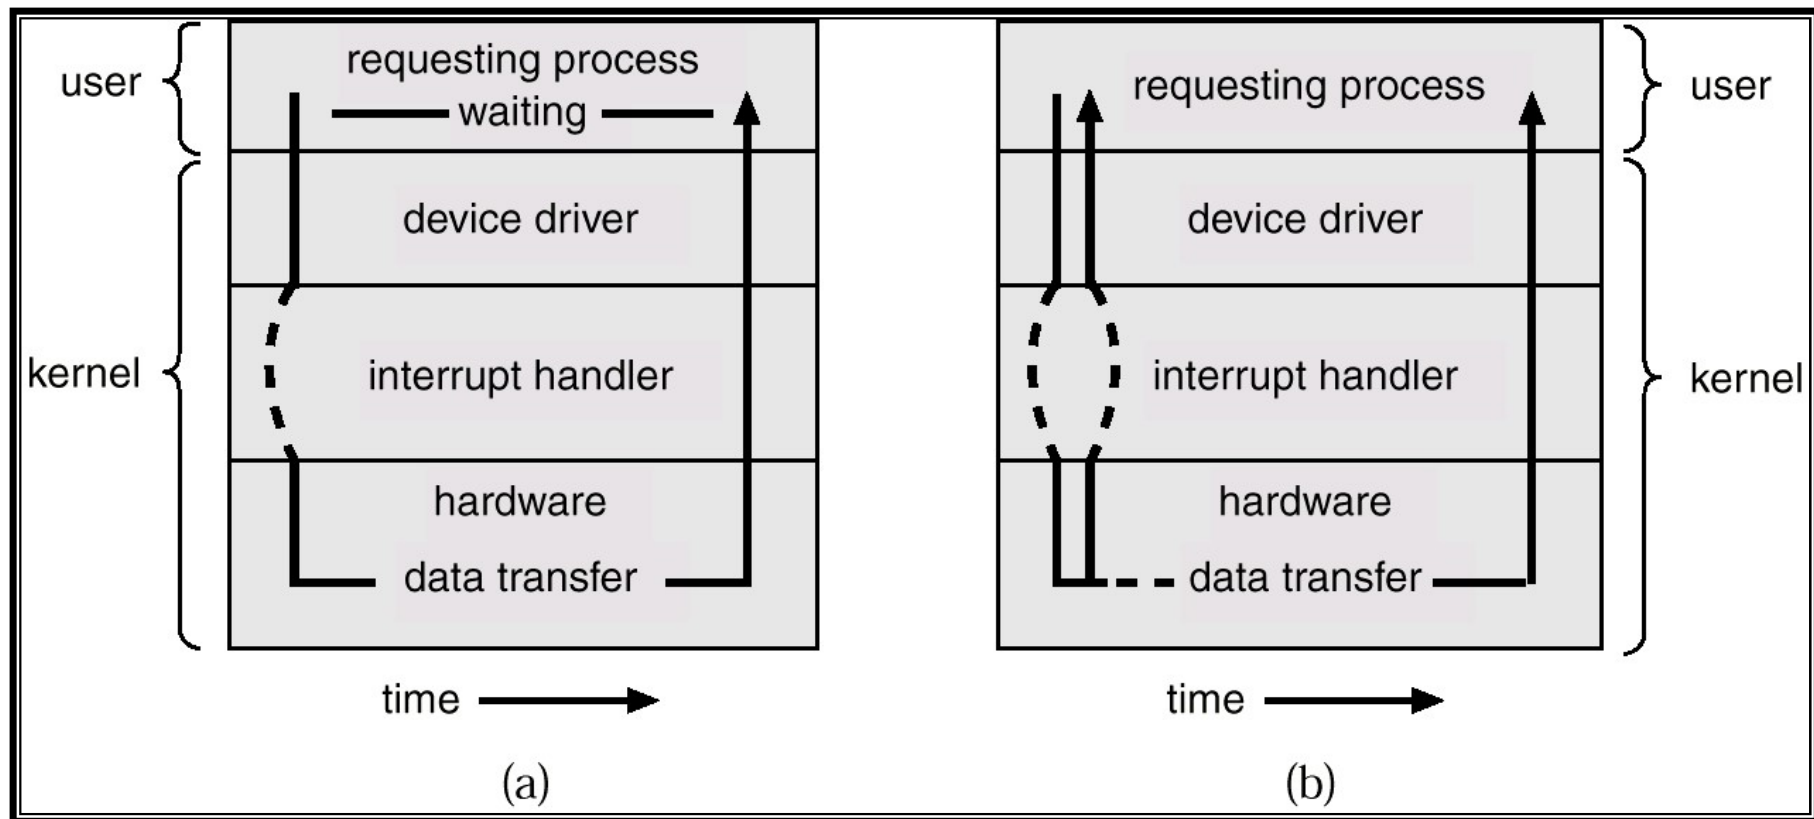
# Interrupt Timeline



# Two I/O Methods

## Synchronous I/O

## Asynchronous I/O

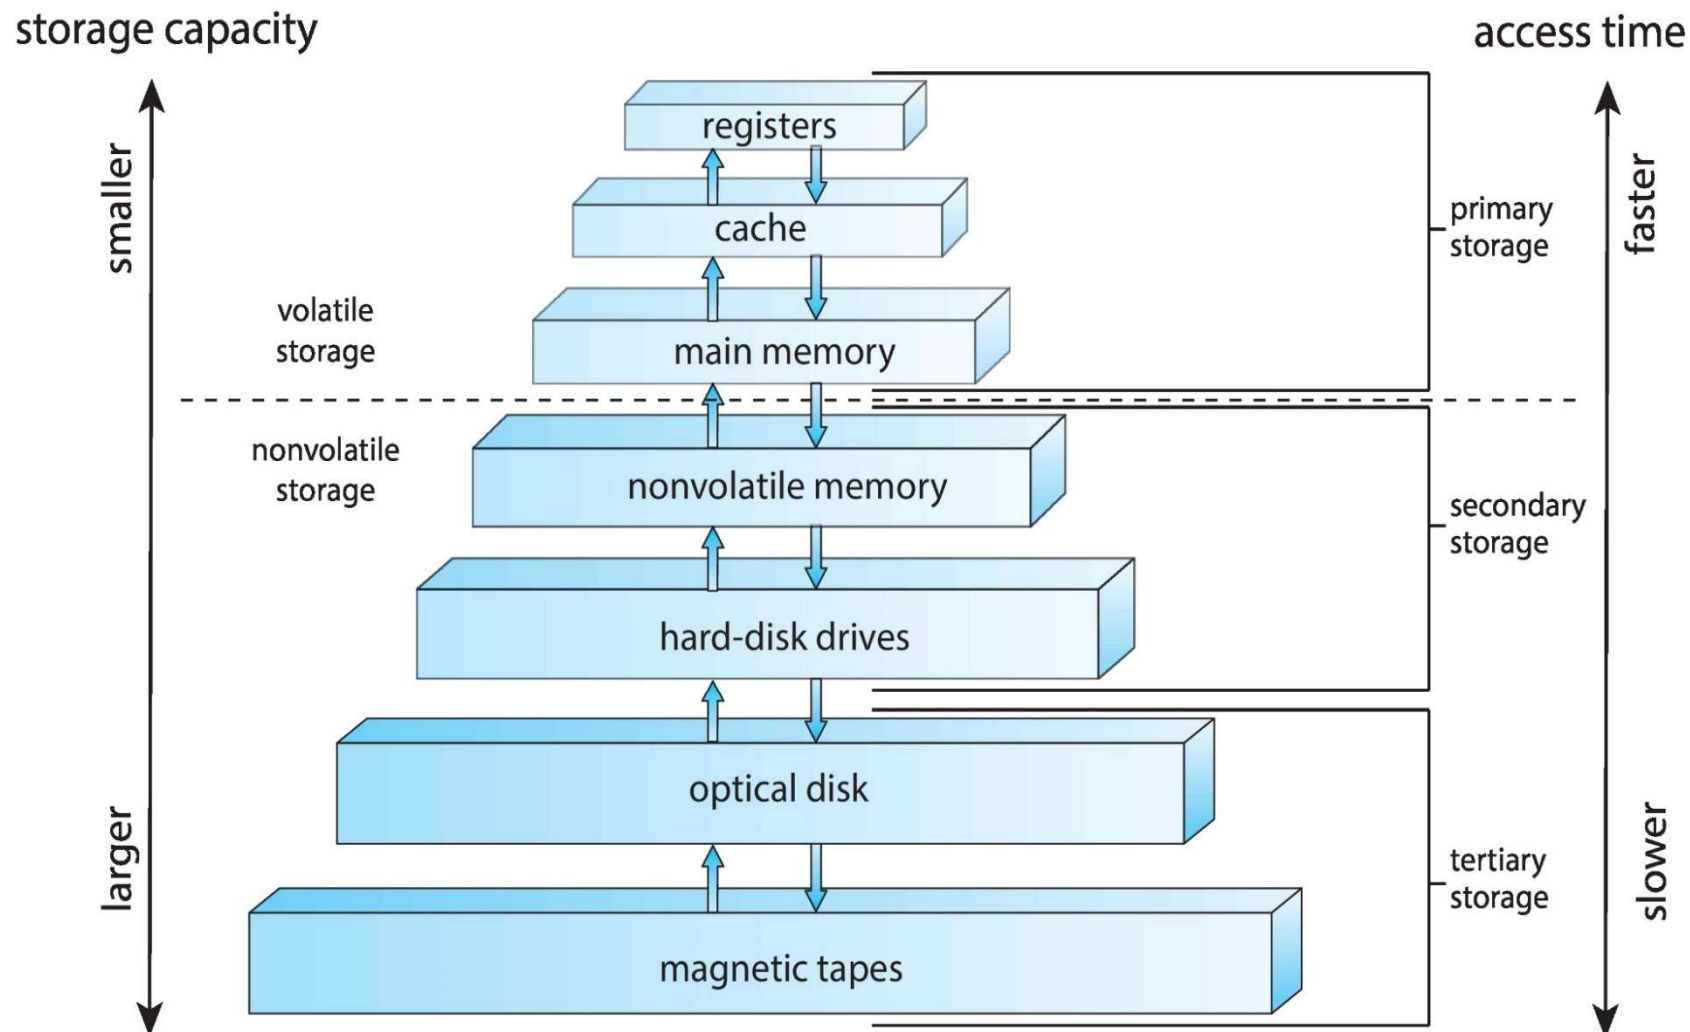


# Direct Memory Access (DMA)

- Who moves the data between memory and device buffer?
  - CPU? device controller?
- DMA
  - Device controller transfers blocks of data from local buffer directly to main memory **without CPU intervention**
  - Raise interrupt when DMA transfer is done

# Storage Structure

- Storage systems organized in **hierarchy**
  - The layers differ in **Sizes, Speed, Cost, Volatility...**



# Storage Structure



- Primary Storage
  - Main memory –the only large storage media that the CPU can access **directly** (e.g., via load/store instructions)
- Secondary storage – extension of main memory that provides large **nonvolatile** storage capacity
  - Magnetic disks (HDD)
    - metal or glass platters covered with magnetic recording material
    - disk surface is logically divided into **tracks/sectors**
    - the most common secondary-storage device
  - Nonvolatile Memory (NVM)
    - SSD (Solid State Drive)
    - Becoming more popular as capacity increases and price drops
    - *Various technologies: **flash memory**, PCM, RRAM, STT-MRAM*
- Tertiary storage
  - Usually for storing backup copies

# Caching

- *Caching* – information in use copied from slower to faster storage temporarily
  - main memory can be viewed as a *cache* for secondary storage
- Faster storage (cache) checked first to determine if information is there
  - If it is, information used directly from the cache (fast)
  - If not, data copied to cache and used there
- Cache is smaller than the storage being cached
  - Cache management is an important design problem
  - Cache size and replacement policy
- Performed at **many levels** in a computer system (in hardware, operating system, software)


# Performance of Various Levels of Storage

| Level                     | 1                                      | 2                             | 3                | 4                | 5                |
|---------------------------|--|-------------------------------|------------------|------------------|------------------|
| Name                      | registers                              | cache                         | main memory      | solid state disk | magnetic disk    |
| Typical size              | < 1 KB                                 | < 16MB                        | < 64GB           | < 1 TB           | < 10 TB          |
| Implementation technology | custom memory with multiple ports CMOS | on-chip or off-chip CMOS SRAM | <b>DRAM</b>      | flash memory     | magnetic disk    |
| Access time (ns)          | 0.25 - 0.5                             | 0.5 - 25                      | 80 - 250         | 25,000 - 50,000  | 5,000,000        |
| Bandwidth (MB/sec)        | 20,000 - 100,000                       | 5,000 - 10,000                | 1,000 - 5,000    | 500              | 20 - 150         |
| Managed by                | compiler                               | hardware                      | operating system | operating system | operating system |
| Backed by                 | cache                                  | main memory                   | disk             | disk             | disk or tape     |

Movement between levels of storage hierarchy can be **explicit** or **implicit** (to software)

- SW-managed caching (explicit) : register, memory, disk
- HW-managed caching (implicit): cache

# Migration of Data $A$ from Disk to Register

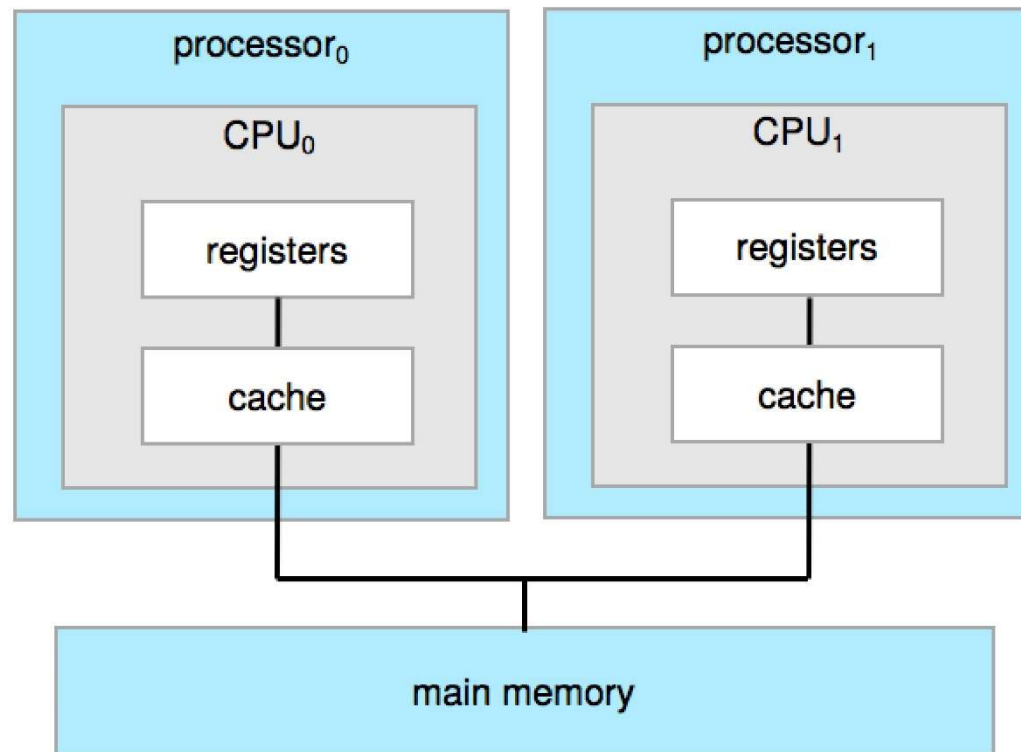




# Computer-System Architecture


- Many systems use **general-purpose processors** (from PCs to mainframes)
  - Some systems have **special-purpose processors** as well
- **Multi-processor systems** growing in use and importance
  - Also known as **parallel systems**, **tightly-coupled systems**
  - Advantages
    1. **Increased throughput**
    2. **Increased reliability** – graceful degradation or fault tolerance
    3. **Economy of scale**
      - Cheaper than multiple single-processor systems due to sharing of resources (e.g., memory, secondary storage, IO devices...)
  - Types
    1. **Asymmetric Multiprocessing**
      - Not all the processors are treated equally
    2. **Symmetric Multiprocessing**

# Symmetric Multiprocessing Architecture



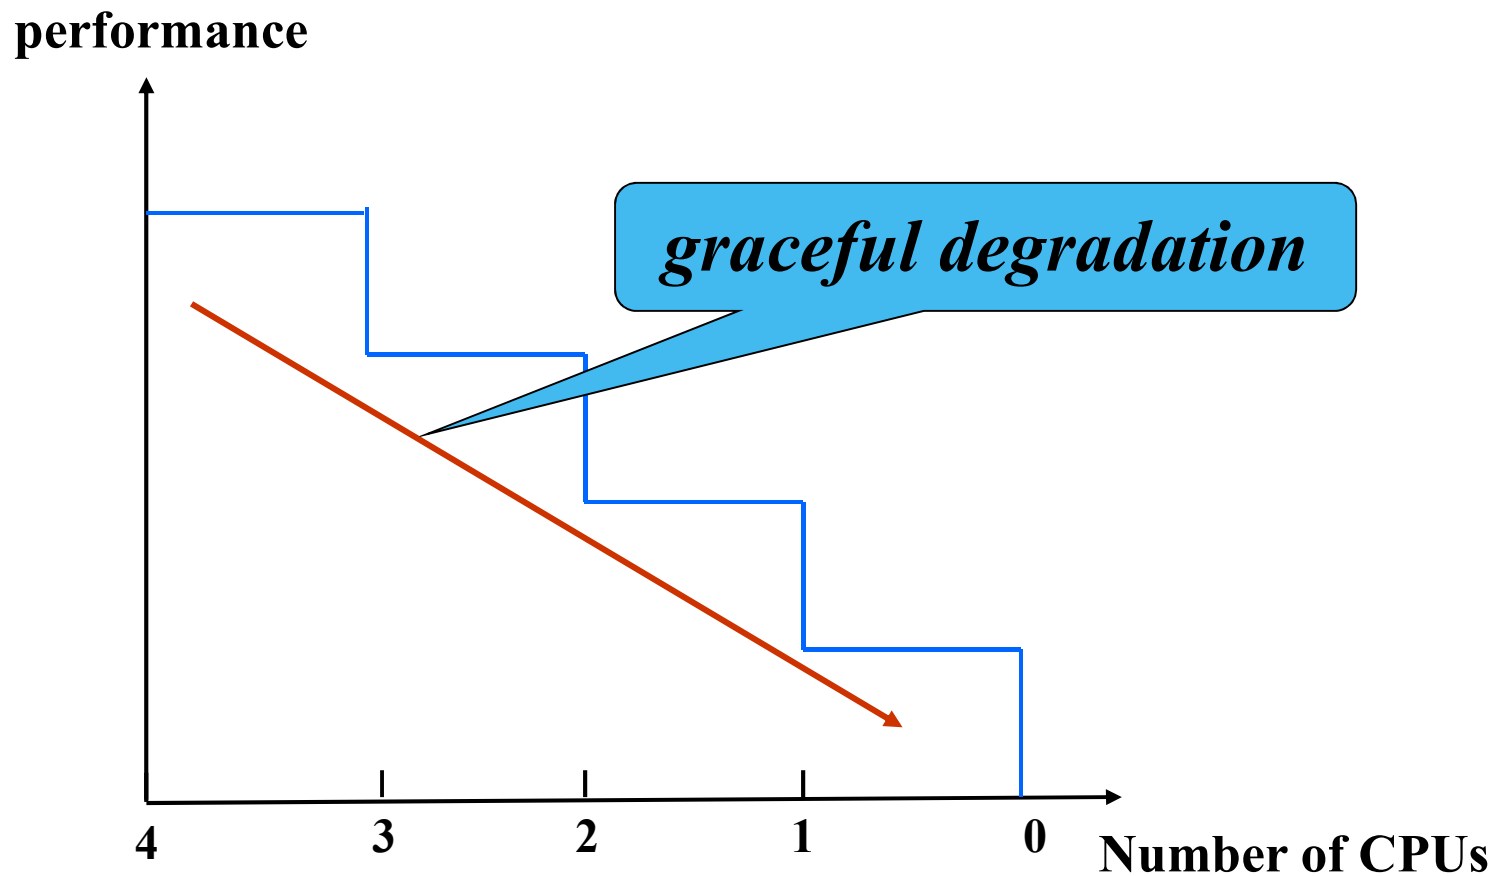
# Multicore Systems

- **Multicore**: multiple **compute cores** on a single chip
  - More efficient than multiple single-core chips
    - On-chip communication is faster than cross-chip communication



**A dual-core design with two cores on the same chip**

# Graceful Degradation




*suppose a four-core system*

# Non-Uniform Memory Access (NUMA)

- Adding too many CPUs on multiprocessor systems will cause
  - contention for **system bus**
    - Since each CPU needs to access main memory
  - Performance begins to degrade
- Alternative approach: **NUMA**
  - Provide each CPU with its own local memory
    - Accessed via a local bus (next slide)
  - Increasingly popular on servers and high-end computing systems

# Non-Uniform Memory Access (NUMA)



- Problem: memory *access time* is not the same
  - remote mem. access latency > local mem. access latency
  - OS need to manage tasks/data carefully
    - Co-locates a task and its data

# Clustered Systems


- A *clustered system*
  - Multiple computers connected via a *local area network (LAN)* or a *faster interconnection network*.
- The key of a clustered system is *high availability*
  - *Fault tolerant*: suffer a failure of any single component and continue operation
    - *Graceful degradation*: continue providing service *proportional* to the level of surviving hardware

# General Structure of a Clustered System

## An example

\*HP C7000 G3 BladeSystem

-15 nodes, 320 cores, 10 Gb Ethernet



<https://ctoservers.co.uk/hp-c7000-g3-bladesystem-w15-x-bl460c-g8--320-cpu-cores1tb-ram-10gbe---cad-cfd--lsdyna--ansys-2101-p.asp>



# Clustered Systems

- *Asymmetric Clustering*

- Some machines are in *hot-standby mode* while the others are running applications.
- A hot-standby machine (i.e., does nothing but) monitors the other machines and becomes active if one server fails.


- *Symmetric Clustering*

- Machines run applications while monitoring each other
  - more efficient as it uses all available hardware

# Multiprogramming & Timesharing

- **Multiprogramming** is needed for **efficiency**
  - a single job cannot keep CPU and I/O devices busy at all times
    - E.g., CPU idle if the job is waiting for IO
  - Multiprogramming organizes jobs (code and data) so CPU always has one to execute
    - **increase CPU utilization**
  - A subset of total jobs in system is kept in memory
    - When one in-memory job has to wait (for I/O for example), OS switches to another job in the memory
  - Other jobs are kept in the **on-disk job pool**
    - A job is selected and put in memory via **job scheduling**

# Memory Layout for a Multiprogrammed System



# Multiprogramming & Timesharing

- **Timesharing (multitasking)** is logical **extension** to multiprogramming in which CPU switches jobs so **frequently** that users can interact with each job while it is running, creating **interactive** computing
  - **Response time** should be  $< 1$  second
  - If several jobs ready to run at the same time  $\Rightarrow$  **CPU scheduling**
    - differences with job scheduling?
  - If processes don't fit in memory, **swapping** moves them in and out to run
  - **Virtual memory** allows execution of processes not completely in memory
    - Allows a program that is larger than the physical memory to run

# Operating System Operations

- An OS must prevent user jobs/processes from damaging the normal operations of the OS
- **Dual-mode operation** allows OS to protect itself
  - **Processor** provides **user mode** and **kernel mode**
    - User processes run in user mode (non-privileged mode)
    - OS runs in kernel mode (privileged mode)
    - Some instructions designated as **privileged**, can only be executed in kernel mode
      - E.g. I/O control, timer, special registers, interrupt related instructions...
    - a **mode bit** is provided
      - to distinguish when system is running user code or kernel code
  - A **system call** changes the mode to the **kernel mode**
    - return from the system call resets the mode to **user mode**

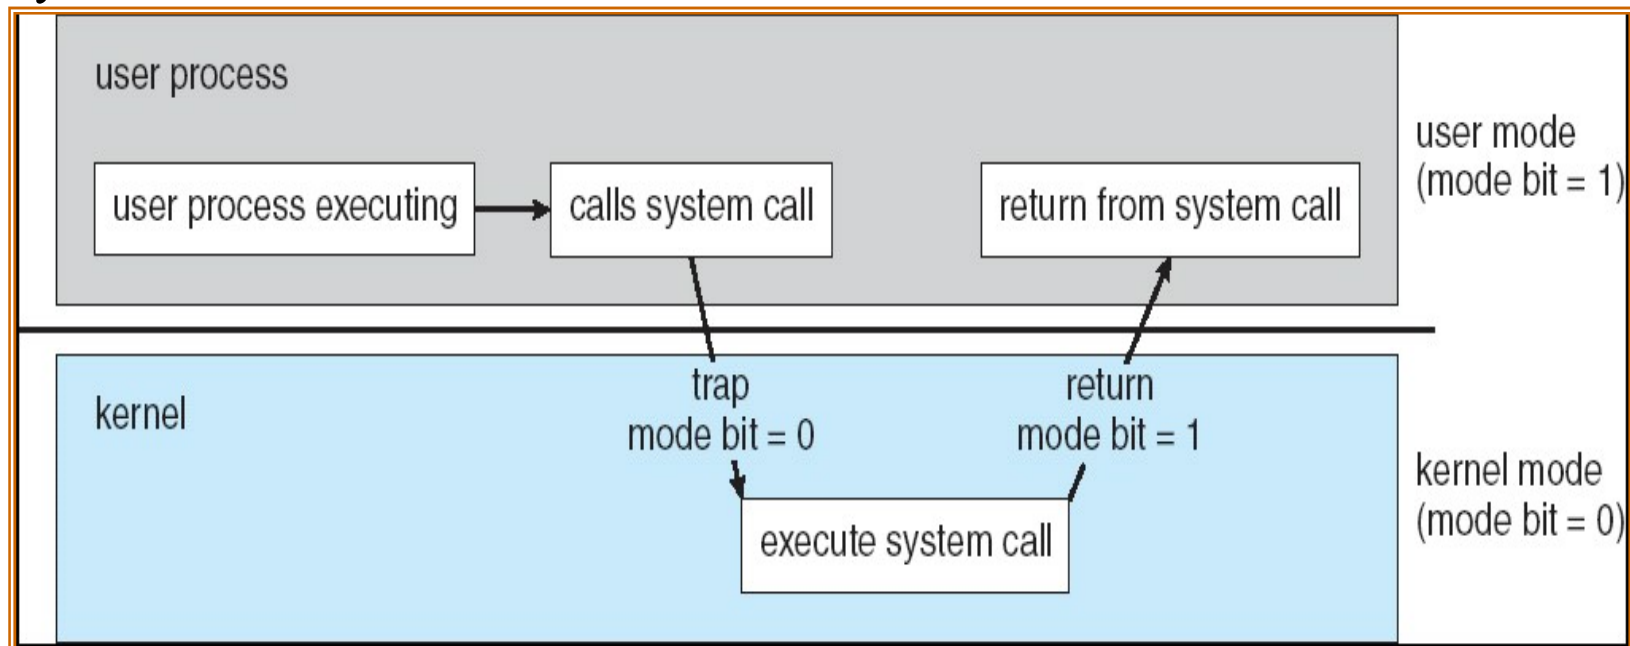
# Transition from User to Kernel Mode

At system boot time, the HW starts in the kernel mode.

- after system initialization, the OS run applications in user mode.


When a user process needs an OS service, it issues an **system call**  
e.g., read a file, send a packet...

System call flow---



# Transition from User to Kernel Mode

- Modes are provided by HW (CPU)
  - Intel's 8088 CPU has only a single mode
    - The OS (e.g., MS-DOS) can be crashed by a user program
  - Modern x86 has (more than) 4 modes/rings



# Timers

- Conceptually, user processes run **on top of** the OS
  - *In fact, they run on the CPU directly!!!*
- How to ensure that the OS maintains control over the CPU?
- Prevent an infinite loop in a process from hogging resources
  - By setting up a timeout period for the process
  - Done by **timers**
    - Operating system decrements the time count for the process periodically (whenever a **timer interrupt** is generated)
    - When the count becomes zero (i.e., the current process has used up its time period)
      - hand the CPU to another process



# Process Management

- A process is a program **in execution**. It is a unit of work within the system.
  - Program is a *passive entity*, process is an *active entity*.
- Process needs resources to accomplish its task
  - CPU, memory, I/O, files
- Process termination requires reclaiming reusable resources
- **Single-threaded** process has **one program counter** specifying location of **next instruction** to execute
  - Process executes instructions sequentially, one at a time, until completion
- **Multi-threaded** process has multiple program counters
  - One program counter per thread
- Typically, a system has many processes running concurrently on one or more CPUs
  - Concurrency by multiplexing the CPUs among the processes/threads

# Process Management Activities



- Activities related to process management
  - Creating and deleting processes
  - Suspending and resuming processes
  - Providing mechanisms for process synchronization
  - Providing mechanisms for process communication
  - Providing mechanisms for deadlock handling
- We will discuss the above topics later

# Memory Management

- **Data** have to be put in memory before being processed
- **Instructions** have to be put in memory before being executed
- Memory management determines **what is in memory**
  - Goal: optimizing CPU utilization and computer's response to users
- Memory management activities
  - Allocating and deallocating memory space as needed
  - Keeping track of **which parts** of memory are currently being used and **by whom**
  - Deciding **which processes** (or parts thereof) **and data** to move into and out of memory

# Storage Management



- OS provides uniform, logical view of information storage
  - Different types of storage devices (e.g., solid-state drives, hard disk drives, tape drives...)
    - **Varying properties** include access speed, capacity, data-transfer rate, access method (sequential or random)
  - **Abstracts physical properties** to **logical storage unit**  
- **file**
  - OS stores files into the storage devices

# Storage Management

- **File-System Management** (provide file interface)
  - File content is determined by its creator
    - Free format – e.g. text files...
    - Fixed format – e.g. executable files...
  - Files usually organized into **directories**
  - **Access control** on most systems to determine **who** can access **what**
  - OS activities include
    - Creating and deleting files and directories
    - Primitives to manipulate files and dirs
      - read/write/append files, set/get file status, set/get file permissions

# Storage Management

- **Mass-Storage Management** (focus on efficient operation)
  - Usually, disks
  - Mass storage devices are used to store
    - data that do not fit in main memory, or
    - data that must be kept for a “long” period of time
  - OS activities
    - Free-space management (bitmap? free list?)
    - Storage allocation
    - Disk scheduling
  - Critical to system performance
    - Because disk is usually the performance bottleneck
    - Latency of a computer operation usually depends on disk subsystem and its algorithms

# I/O Subsystem

- One purpose of OS is to hide peculiarities of hardware devices from the user
- I/O subsystem
  - Drivers
    - Manages specific hardware devices
    - In an OS, only drivers are device-specific
      - Have knowledge about its own device
    - OS provides a general **driver interface** to allow easy cooperation of OS and its drivers
  - **Memory management** of I/O including
    - **buffering** (storing data temporarily while it is being transferred)
    - **caching** (storing parts of data in faster storage for performance)
- Description here is also valid for storage devices
  - Disks are IO devices

# Protection and Security



- **Protection** – mechanisms for controlling **access** (of processes or users) **to resources**
  - Distinguish between authorized and unauthorized usage
- **Security** – defense of the system against internal and external attacks
  - Huge range, including denial-of-service, viruses...




# Protection and Security

- Systems generally first distinguish among **users**, to determine who can do what
  - Each user has a user identifier (**user ID**)
  - User IDs are associated with all files & processes to determine access control
    - **Process owners** and **file owners** are all users...
  - Group identifier (**group ID**) allows set of users to be defined and managed, also associated with each process or file
  - **Privilege escalation** allows user to change to effective ID with more rights *temporarily*
    - **Setuid** in UNIX (see next slide)
      - Causes a program to run with a **user ID of the file owner**, rather than the user's ID
        - » E.g., change your own password, which requires the root privilege to update the password file.
    - Might become a vulnerability if the program is attacked...

# Setuid in Linux

```
cissoll> ls -al passwd  
-r-sr-sr-x 1 root sys 27228 Aug 17 2007 passwd*  
cissoll>
```

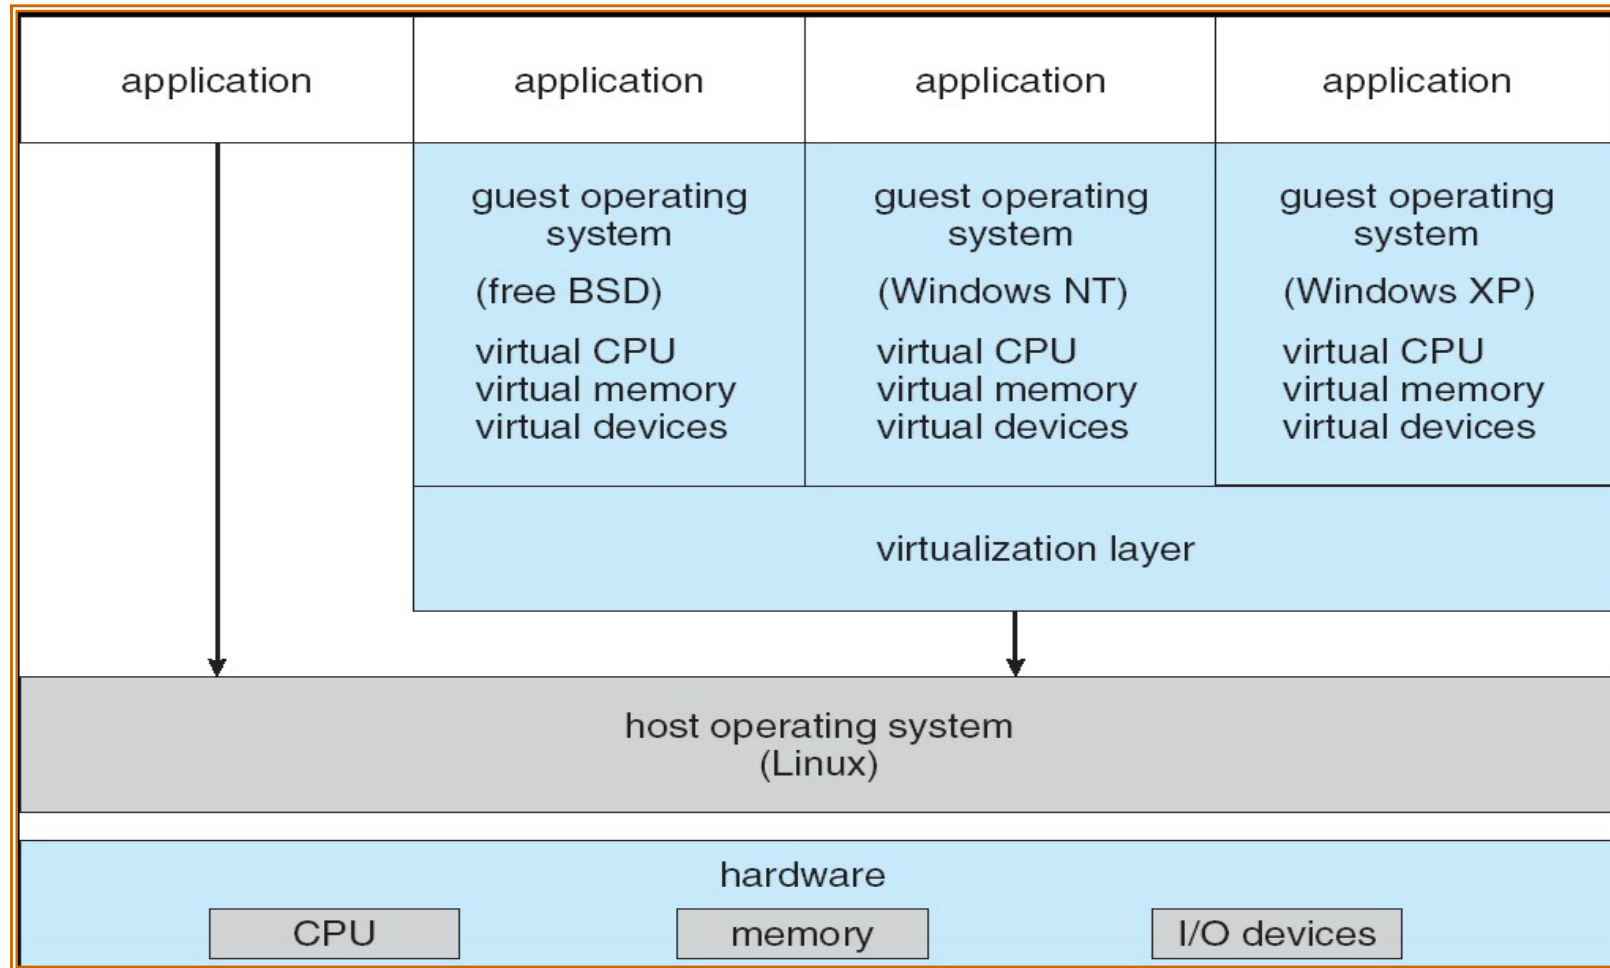


# Virtualization

- **Virtualization**

- Abstract **the hardware of a single computer** into several different execution environments
- Create an illusion that each environment is running on its own private computer
  - An OS can be run in each environment
- Allows OSes to run on other OSes

# Virtualization



*host* operating system vs. *guest* operating system

# Virtualization

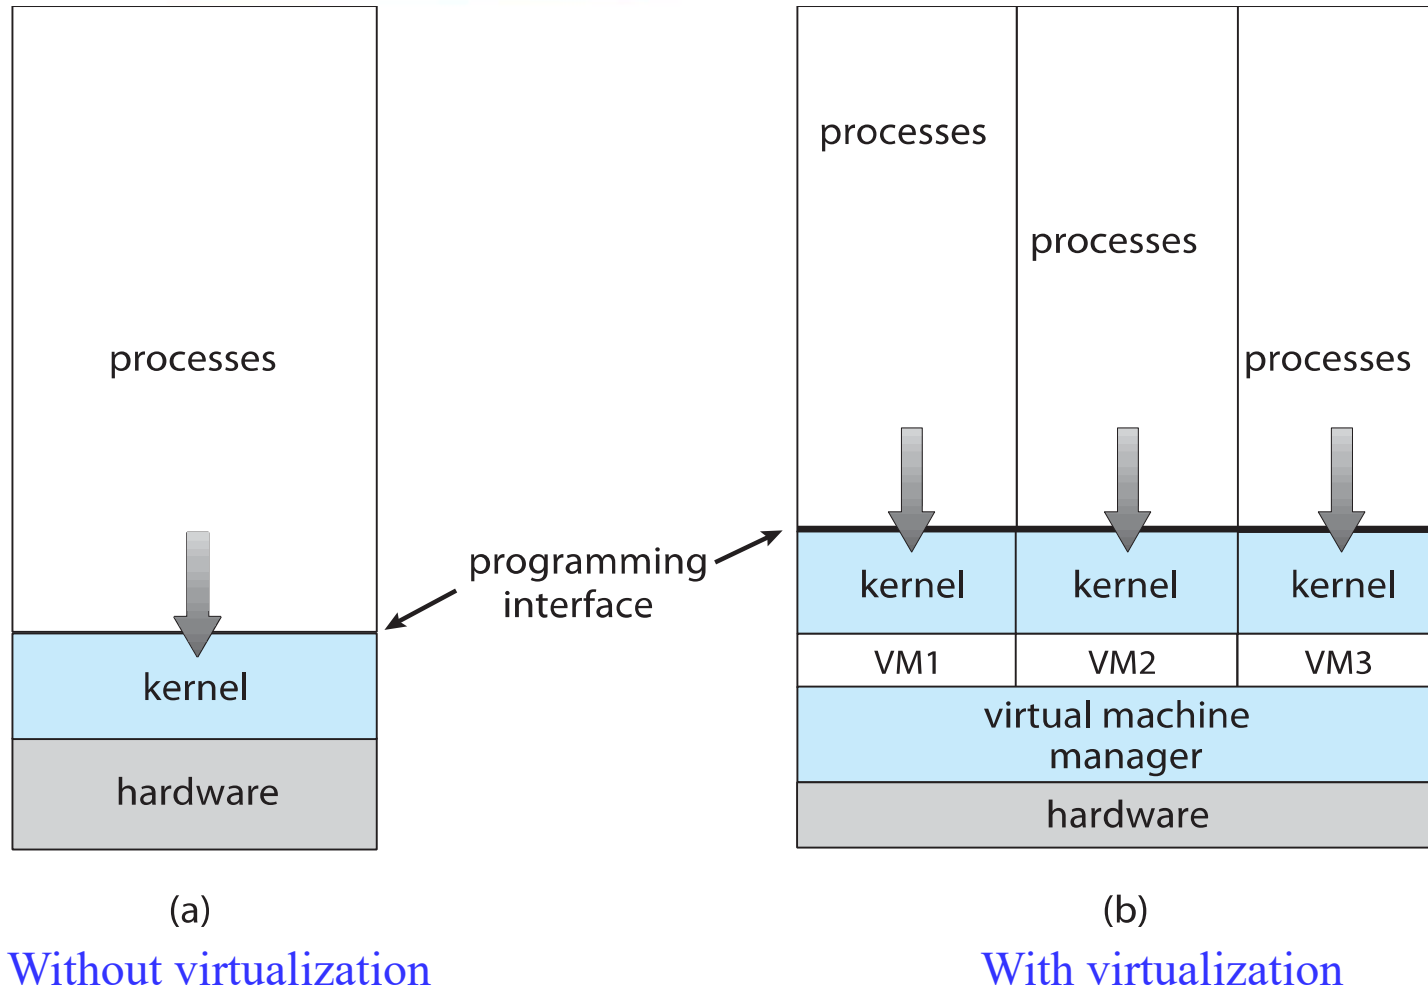
- **Emulation** –

- Simulating computer hardware in software
  - Used when source CPU type **different** from target type
- For Example
  - We want to run the PowerPC applications on the Intel x86
  - Every instruction of PowerPC must be **translated** to the instructions on Intel x86
- Slow...

# Virtualization (Cont.)

- **Virtualization** –
  - An OS **natively compiled for a CPU**, running on another OS **also natively compiled to that CPU**
  - Example
    - Run Linux for x86 on the Windows 10 (also for x86)
    - Linux: **guest OS**
    - Windows 10: **host OS**
  - **VMM (Virtual Machine Manager)** provides virtualization services
    - Also called **Hypervisor**
  - VMM can also run **natively** without the host OS
    - Example: VMware ESX and Citrix XenServer

# VMM without Host OS



*VMM directly runs on hardware (do not need host OS)*


# Kernel Data Structures

- Lists, Stacks, and Queues
- Trees
- Hash Tables
- Linux data structures defined in *include* files
  - e.g., `<linux/list.h>`, `<linux/kfifo.h>`,  
`<linux/rbtree.h>`




# List


- *Singly linked list*



- *Doubly linked list*



- *Circular linked list*




# Stack and Queue

- **Stack:** last in, first out (LIFO)
  - Insert an item: **push**
  - Remove an item: **pop**
- **Queue:** first in, first out (FIFO)

# Tree


- **General tree**
  - A parent may have an unlimited number of children
- **Binary tree**
  - A parent may have at most two children
- **Binary search tree**
  - Two children, left  $\leq$  right
  - Search performance is  $O(n)$
  - **Balanced binary search tree**
    - Search performance is  $O(\lg n)$



**Binary search tree**

# Hashing and Bitmap

- **Hash function** can create a **hash map**



- **Bitmap** – string of  $n$  binary digits representing the status of  $n$  items

# Computing Environments



- Traditional Computing
- Mobile Computing
- Client-Server Computing
- Peer-to-Peer Computing
- Cloud Computing
- Real-Time Embedded Systems

# Computing Environments



- Traditional Computing
  - Office environment
    - PCs connected to a network, **terminals** attached to **mainframe or minicomputers** providing **batch and timesharing** functionalities
    - Now **portals** allowing remote systems to access internal resources
  - Home networks
    - Used to be single system, then modems
    - Now firewalled, networked

# Mobile Computing




- Computing on **handheld devices**
  - Smartphones
  - Tablets
- Issues
  - Limited memory
  - Slower processors
  - I/O constraints: small display screens, small keyboards.....
  - Limited power
- Extra features
  - GPS, accelerometers, and gyroscope
- Leaders are **Apple iOS** and **Google Android**

# Client-Server Computing

**Servers:** serving requests, **Clients:** issuing requests

- ▶ **Application-server** provides an interface to client to request services
- ▶ **File-server** provides interface for clients to store and retrieve files

For clients, dumb terminals are replaced by PCs or smart phones






# Peer-to-Peer Computing

- P2P does not distinguish clients and servers
  - Instead all nodes are considered *peers*
  - May each act as client, server or both
  - Node must join P2P network
- Advantage
  - In client-server system, the server may become a bottleneck.  
In P2P, services can be provided by several nodes
- Example
  - skype

# Peer-to-Peer Computing

- Node must join P2P network
  - Registers its service(s) with central lookup service on network, or
  - Broadcast requests for services and respond to requests for services via *discovery protocol*




# Cloud Computing


- Delivers computing, storage, even apps as a service across a network
- Logical extension of **virtualization** (see next slide)
  - Typically based on virtualization
  - Amazon **EC2** has thousands of servers, millions of **VMs**, PBs of storage available across the Internet, pay based on usage
- Different types
  - **Public cloud** – available via Internet to anyone willing to pay
  - **Private cloud** – run by a company for the company's own use
  - **Hybrid cloud** – includes both public and private cloud components

# Logical Extension of Virtualization


**WWW Server**  
Solaris



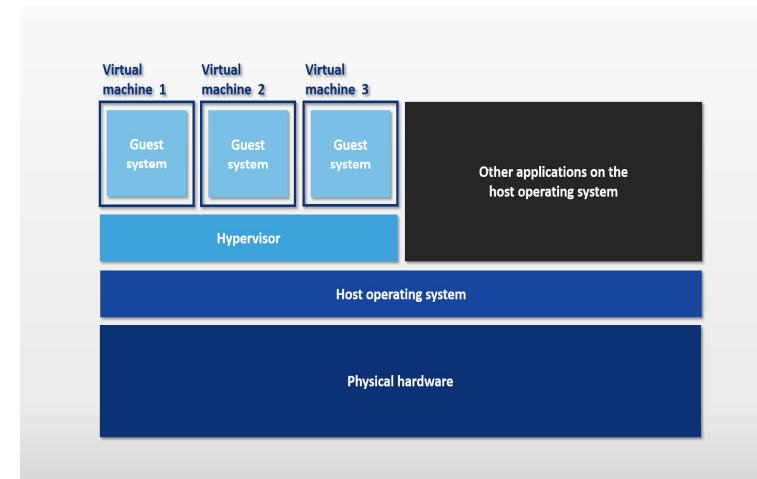
**File Server**  
Win10



**Database Server**  
Linux



**WWW + File + Database  
Servers**




# Cloud Computing



- Categories
  - **Software as a Service (SaaS)** –
    - one or more applications available via cloud (e.g., word processor)
  - **Platform as a Service (PaaS)** –
    - software stack ready for application use via cloud (e.g., a database server)
  - **Infrastructure as a Service (IaaS)** –
    - servers or storage available via cloud (e.g., storage available for backup use)

# Cloud Computing

- Cloud computing environments composed of traditional OSes, plus VMMs, plus cloud management tools
  - Internet connectivity requires security like firewalls
  - Load balancers spread traffic across multiple applications



# Real-Time Embedded Systems



- Embedded computers
  - Most prevalent form of computers
    - Cars, routers, robots, printers,...
  - Tend to have very **specific** tasks
- Different OS implementation schemes
  - No OS
  - A tiny embedded OS
  - A standard OS (e.g., Linux)

# Real-Time Embedded Systems

- Embedded systems sometimes run *real-time operating systems*
  - Processing must be done within a predefined time constraint (i.e., **deadlines**)
  - Otherwise, the system will fail
    - e.g., missile systems



# Open-Source Operating Systems



- Many operating systems are open source
  - e.g., Linux, FreeBSD, OpenSolaris, Illumos,...